

MedIQ: Enhancing Remote Healthcare with Generative AI and Workflow Automation

Anand More

(Pimpri Chinchwad College of Engineering and Research, Pune, India
moreanand111011@gmail.com)

Archana Chaugule

(Pimpri Chinchwad College of Engineering and Research, Pune, India
archana.chaugule@pccoer.in)

Sayyoni Parate

(Pimpri Chinchwad College of Engineering and Research, Pune, India
sayyoniparate@gmail.com)

Sujal Shahare

(Pimpri Chinchwad College of Engineering and Research, Pune, India
sujalshahare2015@gmail.com)

Abstract: As the demand for accessible healthcare grows, telemedicine must evolve beyond simple video calls to actively support both doctors and patients. To bridge this gap, we developed AidME—a Generative AI-powered telehealth ecosystem designed to make remote consultations more engaging and clinically efficient. At its core, AidME facilitates seamless, low-latency video communication using WebRTC. However, its true value lies in its role as an automated clinical assistant: as the consultation takes place, OpenAI Whisper transcribes the conversation in real time. This dialogue is then intelligently distilled by a LLaMA-driven Retrieval-Augmented Generation (RAG) framework, which securely archives the medical history and instantly drafts prescriptions for the doctor’s review. To further support medical professionals, AidME utilizes TensorFlow-based predictive models to analyze patient symptoms and offer early diagnostic insights. We built this unified workflow on a scalable Node.js–React architecture, powered by Kubernetes-orchestrated microservices to guarantee high reliability. Recognizing the critical importance of patient privacy, the platform ensures strict HIPAA and GDPR compliance through robust encryption and continuous system observability. By seamlessly blending real-time communication with automated documentation and AI diagnostics, AidME reduces administrative burnout for clinicians and delivers a smarter, more personalized healthcare experience.

Keywords: Generative AI, LLaMA 3.2 Transformer, Retrieval-Augmented Generation, Natural Language Processing, Summarization Pipeline, TensorFlow / PyTorch Models, Disease Prediction Model, Medical Knowledge Graph, Model Quantization and Pruning, Hybrid Chatbot, Real-Time AI Inference

Categories: J.3, I.2.1, I.2.7, C.2.0, I.5.2

DOI: 10.3897/jucs.<SubmissionNumber>

1 Introduction

The rapid expansion of digital healthcare has reshaped how medical services are delivered, allowing patients to seek consultations beyond the limits of traditional hospital visits.

Yet, despite this progress, many telemedicine platforms still face practical shortcomings. Patients often experience delays during consultations, doctors manage fragmented clinical data spread across disconnected systems, and the lack of automation forces medical professionals to spend a considerable amount of time on routine documentation [Sharma et al. 2023]. These limitations reduce the overall effectiveness of telehealth, especially in settings where timely decisions and reliable information flow are crucial.

AidME is developed to address these gaps through a unified, intelligent telehealth ecosystem that brings together real-time communication, automated documentation, and AI-assisted clinical support. The system integrates WebRTC to facilitate stable, low-latency video consultations between doctors and patients, creating an experience that closely mirrors in-person interactions. During each consultation, the audio stream is processed using Whisper-based speech recognition to generate accurate, real-time medical transcripts. Instead of storing long unstructured text, the platform applies a LLaMA-driven Retrieval-Augmented Generation (RAG) pipeline to summarize, organize, and archive the clinical information in a structured vector database. This makes past consultations easy to retrieve, supports continuity of care, and reduces manual record-keeping tasks.

Beyond transcription and documentation, AidME incorporates machine learning models built with TensorFlow to assist doctors in detecting likely health conditions and preparing prescriptions. These predictive tools do not replace clinical judgement; rather, they support faster decision-making by presenting data-driven insights and suggesting relevant medications based on recognized symptoms. Doctors retain full control over final prescriptions, ensuring that AI remains an assistive component rather than an autonomous authority.

The platform is engineered with a scalable Node.js–React architecture, ensuring smooth performance for both web and mobile environments. All services are containerized with Docker, enabling reliable deployment, while Kubernetes orchestration ensures automatic scaling, load balancing, and fault tolerance as consultation demand increases. Strict adherence to globally accepted standards such as HIPAA and GDPR ensures that patient information remains secure through end-to-end encryption, access control policies, and continuous auditing.

By combining these technologies into a single integrated environment, AidME aims to overcome longstanding inefficiencies in telemedicine. The platform reduces administrative burden on healthcare providers, accelerates diagnostic workflows, and enhances patient access to modern, intelligent medical support. Ultimately, AidME represents a step toward a more connected, automated, and patient-centered model of digital healthcare, capable of meeting the rising demand for remote medical services with accuracy, speed, and reliability.

2 Background and Related Work

AI-Driven Telemedicine Systems: AI-enhanced telemedicine has progressed significantly in recent years, aiming to improve diagnostic workflows, automate documentation, and enhance doctor–patient communication. Prior work highlights persistent challenges in telemedicine, including fragmented data flow, limited decision-support automation, and the absence of unified transcription-to-summary pipelines [Sharma et al. 2023]. These limitations motivate the development of integrated, intelligent platforms capable of real-time processing and clinical reasoning.

Advancements in Speech Recognition for Telehealth: Speech-to-text technologies form the backbone of real-time teleconsultation systems. Faster Speech LLaMA introduced a multi-token prediction mechanism that reduces transcription latency for large-scale speech models [Raj et al. 2024]. Whisper-Flamingo further extends speech recognition by integrating visual features, offering significant improvements in noisy or multimodal environments [Rouditchenko et al. 2024]. Despite strong performance, these models impose high computational requirements and are not optimized for deployment in low-latency, resource-constrained telemedicine settings.

LLM-Based Clinical Reasoning and Disease Prediction: Large language models designed for clinical applications have achieved notable success. DRG-LLaMA demonstrated strong predictive accuracy for automated diagnosis-related group classification [Wang et al. 2024]. Health-LLM applied retrieval-augmented generation (RAG) to improve disease prediction using personalized patient histories [Jin et al. 2024]. While these models excel in structured medical tasks, they rely on static datasets and do not support real-time conversational inputs from doctor–patient video sessions.

Predictive Analytics and Public Health Modeling: Predictive modeling plays a significant role in healthcare analytics. Olushola et al. proposed spatiotemporal models for forecasting disease outbreaks, aiding large-scale epidemiological planning [Olushola et al. 2023]. These contributions provide strong population-level insights but do not extend to individualized telehealth workflows or real-time medical decision support.

Secure and Trustworthy Data Management in Healthcare: Secure, tamper-resistant data management remains a critical requirement in modern telehealth systems. Blockchain-based frameworks have been proposed to ensure integrity, traceability, and secure storage of medical documents [Brijwani et al. 2024]. Additionally, LLM-driven data analysis methods have enabled efficient insight extraction from heterogeneous datasets [Sánchez Pérez et al. 2025]. However, these approaches typically operate offline and lack seamless integration with real-time clinical consultation environments.

Summary of Gaps and Motivation for AidME: Existing literature addresses specific elements of the telemedicine pipeline—speech recognition, predictive modeling, clinical reasoning, or secure storage—but lacks a holistic architecture capable of integrating all components in real time. The proposed AidME platform bridges these gaps by combining Whisper-based real-time transcription, LLaMA-powered RAG summarization, TensorFlow-based disease prediction, and blockchain-inspired secure data management into a unified, scalable ecosystem. This integration enables a continuous, automated clinical workflow suitable for next-generation telehealth services.

3 Research Gap

Despite significant advancements in AI and telemedicine, several gaps remain that hinder the deployment of fully automated, real-time clinical support systems in healthcare environments:

1. **Limited Integration of Real-Time Transcription and Summarization:** Existing studies focus either on AI-assisted transcription [Raj et al. 2024, Rouditchenko et al. 2024] or on medical coding and predictive reasoning [Wang et al. 2024, Jin et al. 2024]. Few approaches integrate live doctor–patient dialogue with automated summarization, limiting workflow efficiency and the ability to generate actionable insights during consultations.

2. **Unvalidated Clinical Adaptation:** Models like Faster Speech LLaMA [Raj et al. 2024] and Whisper-Flamingo [Rouditchenko et al. 2024] demonstrate high transcription accuracy and efficiency, but their performance in real clinical settings remains largely untested. The absence of validation under real-world telemedicine conditions limits adoption in healthcare.
3. **Computational and Resource Constraints:** Advanced multimodal systems [Raj et al. 2024, Brijwani et al. 2024] often require high computational resources and complex infrastructure, making them impractical for smaller clinics or edge devices. Current frameworks do not adequately address the need for lightweight, deployable solutions without compromising accuracy.
4. **Incomplete End-to-End Workflows:** While disease prediction and LLM-based reasoning frameworks exist [Wang et al. 2024, Jin et al. 2024, Sánchez Pérez et al. 2025], they are mostly limited to structured patient data or retrospective analytics. There is a lack of systems that unify transcription, summarization, predictive diagnostics, and prescription generation in a single platform.
5. **Data Security and Compliance Challenges:** Secure healthcare documentation frameworks [Brijwani et al. 2024] ensure transparency and immutability, but high latency and computational costs hinder real-time integration with AI pipelines. Moreover, compliance with regulations like HIPAA is rarely addressed in end-to-end telemedicine platforms.
6. **Generalization and Adaptability:** Existing AI models are often trained on curated datasets [Sharma et al. 2023, Al Nazi and Peng 2024] that do not fully capture the variability of real-world patient interactions, accents, or multimodal inputs. This reduces model robustness and limits generalization to diverse clinical contexts.

Current telemedicine solutions either focus on isolated transcription, predictive modeling, or secure storage, without providing an integrated, real-time, and scalable solution. These gaps motivate the development of the AidME platform, which combines Whisper-powered live transcription, LLaMA-driven summarization, AI-based disease prediction, and blockchain-inspired secure data handling into a unified, HIPAA-compliant system capable of supporting real-time clinical decision-making.

4 Proposed System

The proposed AidME framework is designed to provide a secure and intelligent telemedicine solution integrating real-time doctor-patient consultations, AI-driven diagnostics, and automated prescription generation. The system ensures remote accessibility, accurate diagnosis, and efficient management of patient health records.

System Architecture: The architecture, as shown in Figure 1, comprises two main modules: the Doctor Module and the Patient Module. Both users interact through the AidME portal after secure login or registration. The patient can book appointments, upload medical images (such as chest X-rays or eye scans), and access AI-driven chatbot assistance, while doctors can manage appointments, conduct video consultations, and review diagnostic reports.

The architecture employs a WebRTC-based video call system integrated with the Whisper model for real-time transcription, which is stored in MongoDB for memory

management. AI modules include an ML-based image analysis system for detecting eye diseases and pneumonia, and a Retrieval-Augmented Generation (RAG) engine combined with large language models (LLMs) to generate and verify prescriptions automatically.




Figure 1: Proposed System Architecture

Tools & Technologies: The system leverages Node.js, Express.js, and Socket.IO for backend and signaling, React.js for frontend, and MongoDB for database management. AI components utilize Python, TensorFlow, OpenCV, and Hugging Face Transformers for model deployment and inference.

5 Methodology

The proposed AidME telemedicine system integrates real-time video consultations, intelligent prescription generation, and automated medical image analysis to assist both patients and doctors through a unified digital platform. The implementation framework consists of five key stages: user authentication, consultation management, AI-assisted diagnosis, prescription generation, and deployment.

5.1 User Authentication

The process begins with secure login and registration through the AidME portal. Users are categorized as either patients or doctors, each granted role-specific access. Authentication is managed using JSON Web Tokens (JWT) to ensure data privacy and session integrity. Patient data and appointment records are stored in a MongoDB database for fast retrieval and scalability.

5.2 Consultation Management

The platform employs WebRTC-based video conferencing for real-time communication between doctors and patients. Signaling is managed through Socket.IO, ensuring low-latency and reliable peer-to-peer connection establishment. During consultations, the Whisper Automatic Speech Recognition (ASR) model transcribes spoken communication into text, which is securely logged for medical documentation and further analysis.

5.3 AI-Assisted Diagnosis

Patients can upload diagnostic images, including chest X-rays and retinal scans, through the web interface. The images are processed using machine learning models to detect potential medical conditions. A Convolutional Neural Network (CNN) is utilized for eye disease identification, while a deep CNN architecture detects pneumonia from chest X-rays. The models are trained using publicly available datasets and optimized with transfer learning to enhance performance on limited clinical data.

5.4 Prescription Generation

Following consultation and diagnostic evaluation, the system's RAG (Retrieval-Augmented Generation) engine collaborates with a Large Language Model (LLM) to generate AI-assisted prescriptions. The engine retrieves relevant medical context, cross-verifies with the patient's medical history, and formulates a treatment plan. The generated prescriptions are reviewed by the doctor before final approval and are stored in the patient's electronic medical record for continuity of care.

Require: Live Audio Stream A , Patient Vector Database D_H

Ensure: Structured AI Prescription P

```

1: // Phase 1: Real-Time Speech Recognition
2:  $T \leftarrow \text{WhisperASR}(A)$  {Transcribe live audio to clinical text}
3: // Phase 2: Retrieval-Augmented Generation (RAG) Search
4:  $Q_E \leftarrow \text{EmbeddingModel}(T)$  {Convert transcript to vector embeddings}
5:  $C_E \leftarrow \text{SimilaritySearch}(Q_E, D_H, \text{top}_k = 5)$  {Fetch top 5 relevant history records}

6:  $H_{rel} \leftarrow \text{ExtractText}(C_E)$  {Retrieve historical text}
7: // Phase 3: Prompt Construction & Integration
8:  $Instruction \leftarrow$  "Formulate safe clinical prescription guidelines."
9:  $Prompt \leftarrow \text{Concatenate}(Instruction, T, H_{rel})$ 
10:  $Tokens \leftarrow \text{LLaMA\_Tokenizer}(Prompt)$ 
11: // Phase 4: Large Language Model Inference
12: // Low temperature (0.2) ensures strict medical determinism
13:  $Out \leftarrow \text{LLaMA\_Model.generate}(Tokens, \text{temperature} = 0.2, \text{max\_tokens} = 512)$ 

14:  $P_{raw} \leftarrow \text{Decode}(Out)$ 
15: // Phase 5: Formatting and Output
16:  $P \leftarrow \text{StructuredJSONParser}(P_{raw})$  {Format for Doctor Verification}
17: return  $P$ 

```

5.5 Deployment Setup

The AidME system is deployed using a Node.js and Express.js backend, React.js frontend, and MongoDB database for persistent storage. The AI components run on Python-based microservices integrated with TensorFlow and PyTorch frameworks. The deployment environment supports both cloud and local configurations, allowing doctors to access the platform remotely while ensuring compliance with healthcare data protection standards. The system ensures high availability, secure data transmission, and scalability for real-time telehealth operations. This integrated methodology allows AidME to deliver a reliable, intelligent, and patient-centric telemedicine experience, bridging the gap between virtual consultation and automated medical assistance.

6 Results and Discussion

The proposed AidME platform underwent rigorous evaluation to assess its diagnostic efficacy, system responsiveness, and clinical applicability within real-world telemedicine scenarios. Performance was quantified using standard classification metrics (accuracy, precision, recall, F1-score), system latency measurements, and user satisfaction indices.

6.1 Diagnostic Model Performance

The image-based diagnostic modules were trained and validated on publicly available medical imaging datasets comprising retinal scans and chest X-rays. A standard Convolutional Neural Network (CNN) architecture was employed for ocular disease classification, yielding an overall accuracy of 94.2%. Simultaneously, a deep CNN architecture optimized for pulmonary opacity detection achieved a substantial accuracy of 96.8% in identifying pneumonia.




Figure 2: System Performance Metrics Visualization

To mitigate the effects of varying imaging conditions inherent in remote consultations, we leveraged transfer learning and adaptive data augmentation techniques. These strategies significantly enhanced model generalization, ensuring high precision and recall while minimizing false-positive outcomes.

Diagnostic Module	Architecture	Accuracy	Precision	Recall
Ocular Disease	CNN	94.2%	93.8%	94.0%
Pneumonia	Deep CNN	96.8%	96.5%	97.1%

Table 1: Performance Metrics of Diagnostic Modules

6.2 Augmented Prescription Generation (RAG + LLM)

The automated prescription generation module, powered by a Retrieval-Augmented Generation (RAG) framework integrated with a fine-tuned Large Language Model (LLM), was evaluated for clinical validity. By analyzing patient symptomatology and historical medical records, the model generated preliminary prescriptions that were compared against doctor-verified benchmarks. The generated prescriptions achieved a Semantic Textual Similarity (STS) score of 91%. Clinical evaluation by participating physicians confirmed that the AI-generated prescriptions required only nominal modifications prior to final approval, thereby streamlining the consulting workflow and reducing administrative overhead without compromising patient safety.

6.3 System Responsiveness and Communication Efficacy

Seamless real-time communication is critical for effective teleconsultation. The AidME framework utilizes a WebRTC pipeline coupled with Socket.IO signaling to maintain




Figure 3: Diagnostic Performance Validation Charts

robust peer-to-peer connectivity. Empirical testing over varied bandwidth conditions demonstrated an average transmission latency of 150–200 milliseconds, comfortably within the threshold for acceptable natural conversational flow (ITU-T standards).

Furthermore, the integration of the Whisper Automatic Speech Recognition (ASR) model facilitated real-time clinical transcription. The ASR module recorded an average Word Error Rate (WER) corresponding to a 93% transcription accuracy, ensuring reliable digitization of the clinical encounter.

6.4 User Evaluation and Operational Viability

To measure practical viability and user acceptance, a pilot study was conducted involving a cohort of 50 participants (25 medical practitioners and 25 patients). Post-consultation surveys resulted in a 92% overall satisfaction rate. Participants highlighted the intuitive navigation, reduction in clinical waiting times, and the effectiveness of the AI-driven diagnostic and transcription support as primary advantages.

The architecture's asynchronous data handling and dual-deployment capability (cloud/edge) demonstrated high scalability, indicating that AidME is highly suitable for deployment in resource-constrained rural clinics alongside larger hospital networks.

6.5 Machine Learning Prediction Performance

The AidME platform incorporates a machine learning-based prediction module designed to assist healthcare professionals in identifying potential diseases based on patient symptoms and historical data. This module plays a supportive role in clinical decision-making by providing data-driven insights during consultations.

To evaluate the effectiveness of the prediction system, the model was tested using a combination of training and validation datasets derived from publicly available medical records. The evaluation focused on key performance metrics such as accuracy, precision, recall, and F1-score to ensure reliability and consistency in predictions.

The model achieved an overall accuracy ranging between 90% and 94%, demonstrating its capability to generalize across diverse patient inputs. High precision and recall values indicate that the system is effective in minimizing false positives and false negatives, which is critical in medical applications. The use of optimized machine learning techniques and data preprocessing methods contributed to improved prediction stability.

Additionally, the prediction module was tested in real-time consultation scenarios, where it successfully provided relevant suggestions based on live input data. This significantly reduced the time required for preliminary diagnosis and enhanced the efficiency of the consultation workflow.

Although the system shows promising results, the predictions are intended to assist doctors rather than replace clinical expertise. Final decisions are always made by healthcare professionals, ensuring safety and reliability in medical practice.

7 Conclusions and Future Work

7.1 Conclusions

This paper introduces the AidME telemedicine framework, an advanced, intelligent healthcare ecosystem designed to bridge the gap in remote medical accessibility. By intricately amalgamating WebRTC-driven real-time teleconferencing, AI-augmented image diagnostics, and LLM-powered prescription generation via RAG, AidME facilitates a cohesive, secure, and highly efficient doctor-patient interface.

Empirical validations underscore the system's robustness. The diagnostic classifiers achieved predictive accuracies exceeding 94%, while the automated prescription module demonstrated a 91% contextual alignment with physician standards. Coupled with low-latency communication and highly accurate clinical transcription, AidME presents a scalable, cost-effective solution capable of operating efficiently even in environments with minimal hardware infrastructure.

7.2 Future Work

While the current framework demonstrates substantial efficacy, future iterations will focus on expanding the platform's diagnostic inclusivity. Planned enhancements include the integration of multilingual ASR capabilities to serve diverse linguistic demographics and the incorporation of wearable Internet of Things (IoT) sensors for continuous, real-time vital sign monitoring.

Additionally, we aim to implement federated learning paradigms to generate personalized, privacy-preserving treatment recommendations. Rigorous emphasis will be placed on fortifying cryptographic data security protocols, optimizing edge-device inference latency, and undertaking expansive clinical trials to ensure stringent adherence to global digital healthcare policies and ethical AI standards.

Acknowledgements

The authors would like to express their gratitude to the faculty and research mentors of Pimpri Chinchwad College of Engineering and Research for their continuous guidance and technical support. The development of AidME was inspired by the vision of improving healthcare accessibility through innovation in artificial intelligence and telemedicine technologies.

References

- [Al Nazi and Peng 2024] Al Nazi, Z., Peng, W.: “Large language models in healthcare and medical domain: A review”; *Medical Informatics Review* (2024).
- [Brijwani et al. 2024] Brijwani, G. N., Ajmire, P. E., Mohammad Junaid, M. A., Charasia, S. A., Bhende, D. O.: “Revolutionizing healthcare record management through blockchain”; *Healthcare Technology Letters* (2024).
- [Jin et al. 2024] Jin, M., Yu, Q., Shu, D., Zhang, C., Fan, L., Hua, W., Zhang, Y.: “Health-llm: Personalized retrieval-augmented disease prediction system”; *IEEE Transactions on Biomedical Engineering* (2024).
- [Olushola et al. 2023] Olushola, A., Mart, A., Iqbal, M.: “Predictive modelling for disease outbreak prediction”; *International Journal of Epidemiology* (2023).
- [Raj et al. 2024] Raj, D., Keren, G., Jia, J., Mahadeokar, J., Kalinli, O.: “Faster speech llama inference with multi-token prediction”; *Proc. Conf. on Speech Technologies* (2024).
- [Rouditchenko et al. 2024] Rouditchenko, A., Gong, Y., Thomas, S., Karlinsky, L., Kuehne, H., Feris, R., Glass, J.: “Whisperflamingo: Integrating visual features into whisper”; *IEEE Conference on Multimodal Learning* (2024).
- [Sánchez Pérez et al. 2025] Sánchez Pérez, A., Boukhary, A., Papotti, P., Castejon Lozano, L., Elwood, A.: “An llm-based approach for insight generation in data analysis”; *Journal of Intelligent Data Science* (2025).
- [Sharma et al. 2023] Sharma, S., Rawal, R., Shah, D.: “Addressing the challenges of ai-based telemedicine: Best practices and lessons learned”; *Journal of Digital Health Systems* (2023).
- [Wang et al. 2024] Wang, H., Gao, C., Dantona, C., Hull, B., Sun, J.: “Drg-llama: Tuning llama model to predict diagnosis-related groups”; *npj Digital Medicine*, vol. 7, no. 16 (2024).